

NanoVDB `ReadAccessor` Benchmark — OLD vs NEW, `Acc=<0,1,2>` vs `Acc=<0>`, CPU vs GPU

Benchmarks comparing the `NANOVDB_USE_OLD_ACCESSOR` caching behaviour (ON vs OFF) and `ReadAccessor<0,1,2>` (full 3-level cache, default) vs `ReadAccessor<0>` (leaf-only cache) across access patterns, on CPU (single- and multi-threaded) and GPU. Results are collected from five machines: a **Razer laptop**, a **Dell laptop**, a **desktop**, a **workstation**, and a **MacBook Air M3**. All five use the **corrected leaf-sampling methodology** — coordinates are harvested from the active leaf voxels of a narrow-band level set sphere so every access resolves at a leaf node. Both `ReadAccessor<0,1,2>` and `ReadAccessor<0>` are benchmarked under both OLD and NEW compile modes on all machines.

What is being measured

The `ReadAccessor` caches the tree nodes visited on the previous access so a spatially nearby access can skip the traversal from the root.

- **OLD** (`NANOVDB_USE_OLD_ACCESSOR` defined): the `get()` method uses `if constexpr + else` chaining. For `getValue` (operation `LEVEL=0`), the `if constexpr(LEVEL<=0)` branch is taken and the `else` discards all higher-level cache checks — so the 3-level accessor falls straight to a **root traversal** on any leaf-cache miss, effectively behaving like a leaf-only accessor.
- **NEW** (`NANOVDB_NO_OLD_ACCESSOR` defined): the `else` is removed, so all three cache-level checks execute unconditionally. A leaf-cache miss can still hit the **level-1 / level-2** cache instead of restarting at the root.
- **`ReadAccessor<0,1,2>`** (`DefaultReadAccessor`): maintains leaf, lower-internal, and upper-internal node caches. With NEW, the full 3-level check is active. With OLD, the extra levels are stored but never checked on a read — making it functionally equivalent to a larger `ReadAccessor<0>` struct.
- **`ReadAccessor<0>`**: caches only the leaf node. On a leaf miss it always goes straight to root. Unaffected by the OLD/NEW flag (no level-1/2 to check).

Methodology

What the numbers measure

Every reported figure is the **wall-clock time of a single random-access read** into the grid, i.e. one `accessor.getValue(ijk)` call, expressed in **nanoseconds per lookup** (`ns/lookup`). It is a pure *read-path* micro-benchmark: no values are written, no math is done on the result beyond a running sum (kept `volatile` / written to an output array so the compiler cannot elide the work).

- **CPU-1T** — one thread walks the coordinate list serially with a single accessor. This is a **latency** number: the cost of one dependent lookup.
- **CPU-MT** — the same total work split across all hardware threads (`nanovdb::util::forEach` → `tbb::parallel_for`), each grain using its own accessor. This is a **throughput** number: total time ÷ total lookups with all cores busy.
- **GPU** — the coordinate list is uploaded to the device; each thread processes a 32-coord chunk with its own accessor. Timed with CUDA events (kernel time only, excluding the one-time grid upload). Also a **throughput** number.

The only thing that differs between the OLD and NEW columns is the accessor's cache-fallback logic (a compile-time switch); the data, coordinates, and harness are identical.

The test data

A **fog-volume sphere** built with `nanovdb::tools::createFogVolumeSphere<float>(radius=500, voxelSize=1.0, halfWidth=3.0)`. Its measured properties:

Property	Value
Value type	float
Index bounding box	[-500, -500, -500] ... [500, 500, 500] (1001 ³ voxels)
Dense voxel count of bbox	1,003,003,001 (~1.00 B)
Active voxels	523,592,077 (~523.6 M)
Fill ratio (active / dense bbox)	52.2 %
Upper internal nodes (level 2, 32 ³)	8
Lower internal nodes (level 1, 16 ³)	272
Leaf nodes (level 0, 8 ³)	82,608
Grid size in memory	188,484,640 bytes (~180 MB)

So it is a **fully populated solid ball** — the interior is all active voxels (the 52 % fill ratio is just the volume of a sphere inside its cube, $\pi/6 \approx 0.524$). It is *sparse* in the VDB sense (only the ball is stored, not the empty corners), but the region the benchmark actually samples is **100 % active**: all access patterns are confined to a cube well inside radius 500, so every lookup lands on a real leaf and traverses the tree for real — we are **not** measuring empty-space shortcuts, we are measuring genuine node traversal and cache reuse.

How much work, how it is averaged

Point patterns (Sequential / LeafJump / NodeJump / Random)	1,048,576 lookups each
Stencil	262,144 centres × 27 neighbours = 7,077,888 lookups
Sampled region	cube of half-extent 256 (points) / [0, 64] ³ dense cube (stencil), both fully inside the ball
Repetition	each configuration run 7 times; the median is reported (robust to OS jitter / turbo ramp)
Warm-up	GPU does one untimed warm-up launch before the 7 timed trials

Hardware / build

Five machines were benchmarked; results for each are in separate sections below.

	Razer Laptop	Dell laptop	Desktop	Workstation
CPU	Intel Core i7-9750H — 6 cores / 12 threads; TBB <code>parallel_for</code> , 4096-coord grains	32 HW threads; TBB <code>parallel_for</code> , 4096-coord grains	AMD Ryzen 9 9950X — 16 cores / 32 threads; same TBB settings	AMD Ryzen Threadripper PRO 7975WX — 32 cores / 64 threads; same TBB settings
GPU	NVIDIA Quadro RTX 5000 with Max-Q Design (SM 7.5); 32-coord chunks/thread, 128 threads/block	NVIDIA RTX 5000 Ada Generation Laptop GPU (SM 8.9); 32-coord chunks/thread, 128 threads/block	NVIDIA RTX PRO 6000 Blackwell Max-Q Workstation Edition (SM 12.0); same CUDA settings	NVIDIA RTX 6000 Ada Generation (SM 8.9, 49 GB); same CUDA settings
Build	Release (-O3 -DNDEBUG), C++17 / CUDA 17	same	same	same
OLD vs NEW	selected at compile time per target (-DNANOVDB_USE_OLD_ACCESSOR vs -DNANOVDB_NO_OLD_ACCESSOR)	same	same	same

Access patterns

Pattern	Stride / shape	Cache behaviour
Sequential	1	leaf (level-0) cache always hot
LeafJump	8 (= leaf dim)	leaf cache cold; level-1 warm (NEW only)
NodeJump	128 (= lower-node dim)	leaf + level-1 cold; level-2 warm (NEW only)
Random	uniform in cube	all caches cold — both fall to root
Stencil	3×3×3 = 27 neighbours per centre, dense sweep	mostly same-leaf hits; boundary neighbours spill to level-1 (NEW rescues)

Results — MacBook Air M3 (Apple M3, 8-core CPU; corrected benchmark, CPU only)

This section uses the corrected leaf-sampling methodology. Coordinates are harvested from the active leaf voxels of a narrow-band level set sphere:
nanovdb::tools::createLevelSetSphere<float>(radius=256, voxelSize=1.0, halfWidth=3.0) .

Property	Value
Grid	narrow-band level set sphere (radius 256, voxel 1.0, halfWidth 3.0)
Active leaf voxels harvested	4,939,794
Leaf nodes (level 0)	26,600
Lower internal nodes (level 1)	80
Upper internal nodes (level 2)	8
Access resolving at leaf (verified)	100 %
CPU	Apple M3, 8 HW threads; TBB parallel_for , 4096-coord grains
GPU	none (no CUDA on Apple silicon) — CPU only
Build	Release (-O3 -DNDEBUG), C++17; each printed figure is the median of 7 trials

CPU single-threaded — ns per access (latency)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>	OLD → NEW <0,1,2>
Sequential	20.98	20.35	19.79	20.04	1.06×
LeafJump	135.47	108.52	53.17	104.53	2.55×
NodeJump	123.07	109.29	90.77	107.62	1.36×
Random	218.39	176.14	185.11	165.49	1.18×
Stencil (ns/lookup)	52.06	47.97	40.44	47.43	1.29×

CPU 8-threaded — ns per access (throughput)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>	OLD → NEW <0,1,2>
Sequential	3.98	3.76	3.82	3.73	1.04×
LeafJump	26.13	21.02	11.63	20.70	2.25×
NodeJump	26.83	20.96	18.36	20.69	1.46×
Random	34.67	30.94	37.92	28.90	0.91×
Stencil (ns/lookup)	11.51	11.36	8.79	10.05	1.31×

Stencil detail — ns per whole 27-neighbour stencil

Mode	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
CPU 1 thread	1405.57	1295.28	1091.88	1280.73
CPU 8 threads	310.72	306.74	237.44	271.33

Takeaways — MacBook Air M3 (corrected benchmark)

- **ReadAccessor<0> is unchanged by OLD → NEW** (within ~1–6 % on every pattern and thread count, i.e. noise). This is the clean control the corrected sampling makes possible — the OLD/NEW flag only touches the multi-level cache chaining in <0,1,2> .
- **<0,1,2> NEW wins exactly where multi-level caching applies: LeafJump 2.55× (1T) / 2.25× (8T)** — the headline — plus NodeJump 1.36–1.46× and Stencil 1.29–1.31×. Sequential is a tie (access stays within one leaf, which both variants cache).
- **Random is the one place <0,1,2> does not help** (0.91× on 8T; the 1.18× on 1T is inflated because random picks among the clustered narrow-band voxels retain some upper-node locality). For pure scatter, <0> remains the better choice — it avoids the extra cache-level checks.

Results — Desktop (AMD Ryzen 9 9950X / RTX PRO 6000 Blackwell, SM 12.0) — corrected benchmark

This section uses the **corrected leaf-sampling benchmark** (same methodology as the MacBook Air M3 and Workstation sections). Coordinates are drawn from the active leaf voxels of a narrow-band level set sphere (createLevelSetSphere<float>(radius=256, voxelSize=1.0, halfWidth=3.0)); every lookup resolves at a leaf node.

Property	Value
Grid	narrow-band level set sphere (radius 256, voxel 1.0, halfWidth 3.0)
Active leaf voxels harvested	4,939,794
Leaf nodes (level 0)	26,600
Lower internal nodes (level 1)	80
Access resolving at leaf	100 %
CPU	AMD Ryzen 9 9950X — 16 cores / 32 threads; TBB parallel_for , 4096-coord grains
GPU	NVIDIA RTX PRO 6000 Blackwell Max-Q Workstation Edition (SM 12.0); 32-coord chunks/thread, 128 threads/block
Access count per pattern	1,048,576 per pattern; stencil centres up to 262,144 (filtered for full 27-tap activity)
Repetition	7 trials, median reported

CPU single-threaded — ns per access (latency)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	1.52	0.91	1.47	0.85
LeafJump	7.25	5.52	4.39	5.58
NodeJump	2.96	2.97	4.01	3.00
Random	21.81	21.45	23.13	20.21
Stencil (ns/lookup)	1.112	1.113	1.141	1.234

CPU 32-threaded — ns per access (throughput)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	0.17	0.17	0.19	0.15
LeafJump	0.46	0.50	0.47	0.51
NodeJump	0.28	0.40	0.37	0.31
Random	1.21	1.12	1.44	1.28
Stencil (ns/lookup)	0.107	0.091	0.096	0.093

GPU — ns per access (throughput)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	0.027	0.026	0.026	0.027
LeafJump	0.074	0.073	0.050	0.073
NodeJump	0.087	0.087	0.066	0.087
Random	0.096	0.096	0.121	0.096
Stencil (ns/lookup)	0.0691	0.0689	0.0630	0.0689

OLD → NEW speedup by accessor type

Pattern	CPU-1T <0,1,2>	CPU-1T <0>	CPU-MT <0,1,2>	CPU-MT <0>	GPU <0,1,2>	GPU <0>
Sequential	1.03×	1.07×	0.89×	1.13×	1.04×	0.96×
LeafJump	1.65×	0.99×	0.98×	0.98×	1.48×	1.00×
NodeJump	0.74×	0.99×	0.76×	1.29×	1.32×	1.00×
Random	0.94×	1.06×	0.84×	0.88×	0.79×	1.00×
Stencil	0.97×	0.90×	1.11×	0.98×	1.10×	1.00×

Key: **ReadAccessor<0>** is essentially unchanged by OLD → NEW — confirming the fix only affects the 3-level accessor. NEW <0, 1, 2> wins strongly for LeafJump (**CPU-1T 1.65×**, **GPU 1.48×**). NodeJump shows a notable CPU-1T regression (0.74×) because with the corrected methodology (one real leaf per lower-internal node), the level-1 cache always misses and the extra checks add latency; the GPU benefits from the level-2 hit path (1.32×) despite the same pattern.

Stencil detail — ns per whole 27-neighbour stencil

Platform	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
CPU 1 thread	30.03	30.04	30.81	33.33
CPU 32 threads	2.88	2.45	2.60	2.50
GPU	1.8660	1.8616	1.7002	1.8613

The GPU stencil **improves 10 % with NEW <0, 1, 2>** (1.7002 vs 1.8660 ns/stencil) — in contrast to the Workstation (RTX 6000 Ada, SM 8.9) where it regressed 14 %. On the RTX PRO 6000 Blackwell (SM 12.0) the extra cache-level checks are absorbed and the level-1/2 cache hits across stencil neighbourhood boundaries produce a net gain. **ReadAccessor<0>** (GPU) is unchanged (1.00×), as expected — it has no level-1/2 logic for the flag to affect.

Best-accessor recommendation (NEW mode, this machine)

Pattern	Best choice	Why
Sequential	<0> (0.85 ns 1T)	Leaf cache stays hot; level-1/2 checks add marginal cost
LeafJump	<0, 1, 2> (4.39 ns)	Level-1 cache rescues leaf misses; 1.65× vs <0>
NodeJump	<0> (3.00 ns)	CPU: level-1 always misses, extra checks add latency; GPU: <0, 1, 2> wins
Random	<0> (20.21 ns)	All levels miss; extra checks are pure cost
Stencil (CPU)	<0, 1, 2> (1.112 ns/lkp OLD, 0.096 ns/lkp MT)	MT edge (1.11×); <0> regresses on 1T
Stencil (GPU)	<0, 1, 2> (0.0630 ns/lkp)	Blackwell benefits: 10 % faster than OLD or <0>

Fair platform comparison — best accessor (NEW), full hardware

Workload	CPU-1T	CPU-32T	GPU	GPU vs CPU-32T
Sequential (<0>)	0.85	0.15	0.026	5.8×

Workload	CPU-1T	CPU-32T	GPU	GPU vs CPU-32T
LeafJump (<0,1,2>)	4.39	0.47	0.050	9.4×
Random (<0>)	20.21	1.28	0.096	13.3×
Stencil (<0,1,2>)	1.141	0.096	0.0630	1.5×

Results — Dell laptop (RTX 5000 Ada Generation Laptop GPU, SM 8.9) — corrected benchmark

This section now uses the **corrected leaf-sampling benchmark** (same methodology as the MacBook Air M3, Desktop, and Workstation sections). Coordinates are drawn from the active leaf voxels of a narrow-band level set sphere (`createLevelSetSphere<float>(radius=256, voxelSize=1.0, halfWidth=3.0)`); every lookup resolves at a leaf node. Both `ReadAccessor<0,1,2>` and `ReadAccessor<0>` are benchmarked under OLD and NEW.

Property	Value
Grid	narrow-band level set sphere (radius 256, voxel 1.0, halfWidth 3.0)
Active leaf voxels harvested	4,939,794
Leaf nodes (level 0)	26,600
Lower internal nodes (level 1)	80
Upper internal nodes (level 2)	8
Access resolving at leaf	100 %
CPU	32 HW threads; TBB <code>parallel_for</code> , 4096-coord grains
GPU	NVIDIA RTX 5000 Ada Generation Laptop GPU (SM 8.9); 32-coord chunks/thread, 128 threads/block
Access count per pattern	1,048,576 per pattern; stencil centres up to 262,144 (filtered for full 27-tap activity)
Repetition	7 trials, median reported

CPU single-threaded — ns per access (latency)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	1.63	1.34	1.49	1.34
LeafJump	10.23	9.62	6.48	9.63
NodeJump	5.06	5.07	4.35	5.02
Random	31.12	29.75	32.60	29.90
Stencil (ns/lookup)	2.013	1.945	2.469	2.153

CPU 32-threaded — ns per access (throughput)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	0.43	0.40	0.45	0.43
LeafJump	0.85	0.86	0.69	0.88
NodeJump	0.68	0.66	0.63	0.65
Random	1.61	1.62	1.96	1.66
Stencil (ns/lookup)	0.188	0.200	0.171	0.179

GPU — ns per access (throughput)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	0.065	0.066	0.067	0.066
LeafJump	0.112	0.112	0.097	0.112
NodeJump	0.150	0.149	0.110	0.149
Random	0.221	0.222	0.234	0.222
Stencil (ns/lookup)	0.0845	0.0842	0.0958	0.0845

OLD → NEW speedup by accessor type

Pattern	CPU-1T <0,1,2>	CPU-1T <0>	CPU-MT <0,1,2>	CPU-MT <0>	GPU <0,1,2>	GPU <0>
Sequential	1.09×	1.00×	0.96×	0.93×	0.97×	1.00×
LeafJump	1.58×	1.00×	1.23×	0.98×	1.15×	1.00×
NodeJump	1.16×	1.01×	1.08×	1.02×	1.36×	1.00×
Random	0.95×	1.00×	0.82×	0.98×	0.94×	1.00×
Stencil	0.82×	0.90×	1.10×	1.12×	0.88×	1.00×

Key: **ReadAccessor<0>** is essentially unchanged by OLD → NEW — confirming the fix only touches the 3-level accessor (the ~1.10× on <0> stencil is laptop thermal/turbo noise on the serial run, not a real effect). NEW <0, 1, 2> wins strongly for **LeafJump (CPU-1T 1.58×, CPU-MT 1.23×, GPU 1.15×)** and **NodeJump on GPU (1.36×)**. This machine (SM 8.9 Ada) reproduces the Workstation's GPU-stencil **regression (0.88×)**: with the corrected active-voxel stencil the taps stay in-band, the leaf cache is already hot, and the extra level-1/2 checks add overhead with no benefit.

Stencil detail — ns per whole 27-neighbour stencil

Platform	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
CPU 1 thread	54.36	52.52	66.65	58.12
CPU 32 threads	5.07	5.40	4.62	4.82
GPU	2.2813	2.2728	2.5867	2.2813

On CPU the multi-threaded stencil favours NEW <0, 1, 2> (4.62 vs 5.07 ns/stencil, 1.10×), but the single-threaded stencil regresses (66.65 vs 54.36) — the 1T serial runs are the noisiest on this thermally-constrained laptop. The GPU (SM 8.9 Ada) regresses 12 % with NEW <0, 1, 2>, matching the Workstation's SM 8.9 result; **ReadAccessor<0>** is neutral on GPU (1.00×) as expected.

Best-accessor recommendation (NEW mode, this machine)

Pattern	Best choice	Why
Sequential	<0> (1.34 ns 1T)	Leaf cache stays hot; level-1/2 checks add marginal cost
LeafJump	<0, 1, 2> (6.48 ns)	Level-1 cache rescues leaf misses; 1.58× vs <0>
NodeJump	<0, 1, 2> (4.35 ns; GPU 0.110 vs 0.149)	Level-1/2 cache pays off on both CPU and GPU
Random	<0> (29.90 ns)	All levels miss; extra checks are pure cost
Stencil (CPU)	<0, 1, 2> (0.171 ns/lkp MT)	MT edge (1.10×); 1T noisy
Stencil (GPU)	<0> (0.0845 ns/lkp)	NEW <0, 1, 2> regresses 12 %; leaf cache already hot

Fair platform comparison — best accessor (NEW), full hardware

Workload	CPU-1T	CPU-32T	GPU	GPU vs CPU-32T
Sequential (<0>)	1.34	0.43	0.066	6.5×
LeafJump (<0, 1, 2>)	6.48	0.69	0.097	7.1×

Workload	CPU-1T	CPU-32T	GPU	GPU vs CPU-32T
Random (<0>)	29.90	1.66	0.222	7.5×
Stencil GPU (<0>)	—	—	0.0845	—

Results — Workstation (AMD Ryzen Threadripper PRO 7975WX / RTX 6000 Ada, SM 8.9) — corrected benchmark

This section uses the **corrected leaf-sampling benchmark** (same methodology as the MacBook Air M3 section). Coordinates are drawn from the active leaf voxels of a narrow-band level set sphere (createLevelSetSphere<float>(radius=256, voxelSize=1.0, halfWidth=3.0)); every lookup resolves at a leaf node.

Property	Value
Grid	narrow-band level set sphere (radius 256, voxel 1.0, halfWidth 3.0)
Active leaf voxels harvested	4,939,794
Leaf nodes (level 0)	26,600
Lower internal nodes (level 1)	80
Access resolving at leaf	100 %
CPU	AMD Ryzen Threadripper PRO 7975WX — 32 cores / 64 threads; TBB parallel_for , 4096-coord grains
GPU	NVIDIA RTX 6000 Ada Generation (SM 8.9, 49 GB); 32-coord chunks/thread, 128 threads/block
Access count per pattern	1,048,576 per pattern; stencil centres up to 262,144 (filtered for full 27-tap activity)
Repetition	7 trials, median reported

CPU single-threaded — ns per access (latency)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	1.59	1.29	1.26	1.10
LeafJump	8.06	6.44	4.41	6.52
NodeJump	5.30	4.32	4.67	4.47
Random	23.09	19.04	27.28	19.00
Stencil (ns/lookup)	1.857	1.561	1.761	1.716

CPU 64-threaded — ns per access (throughput)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	0.22	0.18	0.22	0.18
LeafJump	0.36	0.37	0.31	0.32
NodeJump	0.29	0.27	0.34	0.29
Random	0.74	0.68	0.81	0.67
Stencil (ns/lookup)	0.066	0.077	0.065	0.073

GPU — ns per access (throughput)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	0.027	0.026	0.026	0.026
LeafJump	0.063	0.063	0.045	0.063
NodeJump	0.086	0.086	0.058	0.083

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Random	0.133	0.132	0.147	0.132
Stencil (ns/lookup)	0.0587	0.0586	0.0669	0.0583

OLD → NEW speedup by accessor type

Pattern	CPU-1T <0,1,2>	CPU-1T <0>	CPU-MT <0,1,2>	CPU-MT <0>	GPU <0,1,2>	GPU <0>
Sequential	1.26×	1.17×	1.00×	1.00×	1.04×	1.00×
LeafJump	1.83×	0.99×	1.16×	1.16×	1.40×	1.00×
NodeJump	1.13×	0.97×	0.85×	0.93×	1.48×	1.04×
Random	0.85×	1.00×	0.91×	1.01×	0.90×	1.00×
Stencil	1.05×	0.91×	1.02×	1.05×	0.88×	1.01×

Key: **ReadAccessor<0>** is essentially unchanged by OLD → NEW — confirming that the fix only affects the 3-level accessor. NEW <0,1,2> wins strongly for LeafJump (CPU 1T 1.83×, GPU 1.40×) and NodeJump on GPU (1.48×). Stencil shows only a small CPU win and a slight GPU regression — explained below.

Stencil detail — ns per whole 27-neighbour stencil

Platform	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
CPU 1 thread	50.14	42.14	47.53	46.34
CPU 64 threads	1.79	2.08	1.75	1.97
GPU	1.5859	1.5828	1.8057	1.5742

The GPU stencil **regresses 14 % with NEW <0,1,2>** (1.8057 vs 1.5859 ns/stencil). With the corrected methodology the stencil centres are active interior-band voxels whose full 27-neighbour neighbourhood is also active — meaning all taps stay within the narrow leaf band, the level-0 (leaf) cache is already hot, and the extra level-1/2 checks in NEW add overhead with no benefit. With the corrected stencil the taps stay entirely within the narrow-band leaves so the level-1 cache provides no benefit and the extra checks are pure overhead.

Best-accessor recommendation (NEW mode, this machine)

Pattern	Best choice	Why
Sequential	<0> (1.10 ns 1T)	Leaf cache stays hot; level-1/2 checks add marginal cost
LeafJump	<0,1,2> (4.41 ns)	Level-1 cache rescues leaf misses; 1.83× vs <0>
NodeJump	<0,1,2> (4.67 ns) or <0> (4.47 ns)	Level-2 gives slight edge on GPU; effectively tied on CPU
Random	<0> (19.00 ns)	All levels miss; extra checks are pure cost
Stencil (CPU)	<0,1,2> (1.761 ns/lkp 1T)	Small edge; taps stay in-leaf so benefit is modest
Stencil (GPU)	<0> (0.0583 ns/lkp)	NEW <0,1,2> regresses; leaf cache already hot

Fair platform comparison — best accessor (NEW), full hardware

Workload	CPU-1T	CPU-64T	GPU	GPU vs CPU-64T
Sequential (<0>)	1.10	0.18	0.026	6.9×
LeafJump (<0,1,2>)	4.41	0.31	0.045	6.9×
Random (<0>)	19.00	0.67	0.132	5.1×
Stencil CPU (<0,1,2>)	1.761	0.065	—	—
Stencil GPU (<0>)	—	—	0.0583	~1.1× vs CPU-64T

Results — Razer Laptop (Intel Core i7-9750H / Quadro RTX 5000 Max-Q, SM 7.5) — corrected benchmark

This section uses the **corrected leaf-sampling benchmark** (same methodology as all other machines). Coordinates are drawn from the active leaf voxels of a narrow-band level set sphere (`createLevelSetSphere<float>(radius=256, voxelSize=1.0, halfWidth=3.0)`); every lookup resolves at a leaf node.

Property	Value
Grid	narrow-band level set sphere (radius 256, voxel 1.0, halfWidth 3.0)
Active leaf voxels harvested	4,939,794
Leaf nodes (level 0)	26,600
Lower internal nodes (level 1)	80
Upper internal nodes (level 2)	8
Access resolving at leaf	100 %
CPU	Intel Core i7-9750H — 6 cores / 12 threads; TBB <code>parallel_for</code> , 4096-coord grains
GPU	NVIDIA Quadro RTX 5000 with Max-Q Design (SM 7.5); 32-coord chunks/thread, 128 threads/block
Access count per pattern	1,048,576 per pattern; stencil centres up to 262,144 (filtered for full 27-tap activity)
Repetition	7 trials, median reported

CPU single-threaded — ns per access (latency)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	2.62	2.16	2.37	2.32
LeafJump	10.28	9.87	6.32	9.97
NodeJump	7.43	7.27	6.57	7.19
Random	35.16	33.95	40.56	34.28
Stencil (ns/lookup)	3.192	2.584	2.912	3.952

CPU 12-threaded — ns per access (throughput)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	0.86	0.81	0.87	0.84
LeafJump	2.02	2.00	1.50	2.19
NodeJump	2.16	1.67	1.39	1.55
Random	5.72	5.82	6.82	5.75
Stencil (ns/lookup)	0.815	0.745	0.660	0.793

GPU — ns per access (throughput)

Pattern	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
Sequential	0.775	0.777	0.838	0.790
LeafJump	0.983	0.986	1.449	1.000
NodeJump	0.432	0.434	0.565	0.504
Random	1.509	1.517	1.697	1.493
Stencil (ns/lookup)	0.2162	0.2169	0.2703	0.2686

OLD → NEW speedup by accessor type

Pattern	CPU-1T <0,1,2>	CPU-1T <0>	CPU-MT <0,1,2>	CPU-MT <0>	GPU <0,1,2>	GPU <0>
Sequential	1.11×	0.93×	0.99×	0.96×	0.92×	0.98×
LeafJump	1.63×	0.99×	1.35×	0.91×	0.68×	0.99×
NodeJump	1.13×	1.01×	1.55×	1.08×	0.76×	0.86×
Random	0.87×	0.99×	0.84×	1.01×	0.89×	1.02×
Stencil	1.10×	0.65×	1.23×	0.94×	0.80×	0.81×

Key: **NEW <0,1,2> wins on CPU** for LeafJump (**1.63×** 1T, **1.35×** MT) and NodeJump MT (**1.55×**) — the corrected leaf-based patterns exercise real level-1/2 cache rescues that the NEW accessor provides. **The GPU (SM 7.5 Turing) regresses on all patterns with NEW <0,1,2>** — the extra cache-level checks add overhead that the Turing architecture cannot absorb; this contrasts with newer GPUs (Ada SM 8.9, Blackwell SM 12.0) where LeafJump and NodeJump NEW wins are 1.15–1.48×. **ReadAccessor<0> is largely unaffected by OLD → NEW on CPU** (0.91–1.08×); the stencil 1T result (0.65×) is likely thermal noise on this constrained laptop — the MT result (0.94×) is neutral.

Stencil detail — ns per whole 27-neighbour stencil

Platform	OLD <0,1,2>	OLD <0>	NEW <0,1,2>	NEW <0>
CPU 1 thread	86.18	69.78	78.62	106.70
CPU 12 threads	22.00	20.10	17.82	21.41
GPU	5.8386	5.8551	7.2970	7.2512

The GPU stencil regresses 20 % with NEW <0,1,2> (7.2970 vs 5.8386 ns/stencil) — consistent with all other GPU patterns showing SM 7.5 does not benefit from the extra cache checks with the corrected methodology. CPU multi-threaded stencil improves 19 % with NEW <0,1,2> (17.82 vs 22.00 ns/stencil).

Best-accessor recommendation (NEW mode, this machine)

Pattern	Best choice	Why
Sequential	<0> (2.32 ns 1T)	Leaf cache stays hot; level-1/2 checks add marginal cost
LeafJump	<0,1,2> (6.32 ns)	Level-1 cache rescues leaf misses; 1.63× vs <0>
NodeJump	<0,1,2> (6.57 ns)	Level-1 rescues; CPU wins (GPU: OLD <0,1,2> is faster)
Random	<0> (34.28 ns)	All levels miss; extra checks are pure cost
Stencil (CPU)	<0,1,2> (0.660 ns/lkp MT)	MT wins 1.23×
Stencil (GPU)	OLD <0,1,2> (0.2162 ns/lkp)	SM 7.5 regresses with NEW; use OLD binary for GPU

Fair platform comparison — best accessor (NEW), full hardware

Workload	CPU-1T	CPU-12T	GPU (NEW <0>)	GPU vs CPU-12T
Sequential (<0>)	2.32	0.84	0.790	1.1×
LeafJump (<0,1,2>)	6.32	1.50	1.000	1.5×
Random (<0>)	34.28	5.75	1.493	3.9×
Stencil (<0,1,2> CPU)	2.912	0.660	—	—

Cross-machine comparison (Razer Laptop vs Dell laptop vs Desktop vs Workstation)

Scope: all five machines (Dell laptop, Desktop, Workstation, MacBook Air M3, and Razer Laptop) use both accessor types and the **corrected** leaf-sampling methodology (coordinates from active leaf

voxels of a narrow-band level set sphere). Results are directly comparable across sections.

1. LeafJump is the biggest CPU win on all machines

With coordinates drawn from actual active leaf voxels, LeafJump (one voxel per leaf in storage order) exercises a realistic cache pattern: the level-0 cache always misses, but adjacent leaves share the same lower-internal node — exactly the case NEW $\langle 0, 1, 2 \rangle$ rescues.

Platform	LeafJump CPU-1T OLD $\rightarrow$ NEW $\langle 0, 1, 2 \rangle$	LeafJump GPU OLD $\rightarrow$ NEW $\langle 0, 1, 2 \rangle$
Desktop (RTX PRO 6000 Blackwell)	1.65×	1.48×
Workstation (RTX 6000 Ada)	1.83×	1.40×
Dell laptop (RTX 5000 Ada)	1.58×	1.15×
Razer Laptop (RTX 5000 Turing)	1.63×	0.68×

CPU LeafJump is consistently 1.58–1.83× across all machines. GPU LeafJump gains are architecture- dependent: Ada (SM 8.9) and Blackwell (SM 12.0) benefit (1.15–1.48×), while Turing (SM 7.5) regresses (0.68×) — the extra level-1/2 checks are not absorbed on this older target.

2. GPU NodeJump wins on Ada and Blackwell; regresses on Turing

Platform	NodeJump GPU OLD $\rightarrow$ NEW $\langle 0, 1, 2 \rangle$
Desktop (RTX PRO 6000 Blackwell)	1.32×
Workstation (RTX 6000 Ada)	1.48×
Dell laptop (RTX 5000 Ada)	1.36×
Razer Laptop (RTX 5000 Turing)	0.76×

With one voxel per lower-internal node from the actual grid topology (staying within the same upper node), the level-2 cache hit path in NEW pays off on Ada and Blackwell GPUs. On SM 7.5 Turing the extra checks add overhead with no net benefit (0.76×).

3. Stencil GPU: result is architecture-dependent

Machine	GPU Stencil OLD $\rightarrow$ NEW $\langle 0, 1, 2 \rangle$
Desktop (SM 12.0 Blackwell)	1.10× (improvement)
Workstation (SM 8.9 Ada)	0.88× (regression)
Dell laptop (SM 8.9 Ada)	0.88× (regression)
Razer Laptop (SM 7.5 Turing)	0.80× (regression)

With the corrected stencil (centres where all 27 neighbours are active leaf voxels), taps stay within the narrow band and the level-0 leaf cache is already hot. On **SM 8.9 (Ada)** — both the Workstation (RTX 6000 Ada) and the Dell laptop (RTX 5000 Ada) — the extra level-1/2 checks in NEW add overhead with no net cache benefit $\rightarrow$ **12–14 % regression** (a clean cross-machine reproduction). On **SM 7.5 (Turing)**, the regression is even larger at **20 %** — the oldest architecture absorbs the checks least efficiently. On SM 12.0 (Blackwell), the architectural improvements absorb the extra checks and the level-1/2 hits at neighbourhood boundaries produce a 10 % gain. `ReadAccessor<0>` is neutral on GPU across all machines.

4. CPU MT pattern: consistent LeafJump and stencil wins

Pattern	Desktop (32T) CPU-MT $\langle 0, 1, 2 \rangle$	Workstation (64T) CPU-MT $\langle 0, 1, 2 \rangle$	Dell laptop (32T) CPU-MT $\langle 0, 1, 2 \rangle$	Razer Laptop (12T) CPU-MT $\langle 0, 1, 2 \rangle$
Sequential	0.89×	1.00×	0.96×	0.99×
LeafJump	0.98×	1.16×	1.23×	1.35×
NodeJump	0.76×	0.85×	1.08×	1.55×

Pattern	Desktop (32T) CPU-MT <0,1,2>	Workstation (64T) CPU-MT <0,1,2>	Dell laptop (32T) CPU-MT <0,1,2>	Razer Laptop (12T) CPU-MT <0,1,2>
Random	0.84×	0.91×	0.82×	0.84×
Stencil	1.11×	1.02×	1.10×	1.23×

LeafJump and stencil show consistent MT gains across all machines (1.02–1.35× and 1.10–1.23× respectively). The Razer Laptop shows the strongest MT gains for LeafJump and NodeJump — its 6-core configuration produces less cross-grain interference, letting the level-1/2 cache hits matter more. Random is consistently near or below 1.00× (extra checks, no cache benefit).

5. ReadAccessor<0> unaffected by OLD → NEW on all platforms

Confirmed: <0> OLD → NEW ratios remain 0.97–1.05× across all patterns, platforms, and methodologies. The fix is entirely in the 3-level <0,1,2> accessor.

6. Random always regresses for <0,1,2>; <0> is neutral

For <0,1,2> : all caches miss → 0.85× on CPU-1T, 0.90× on GPU (corrected run). For <0> : no extra checks → 1.00× on all platforms. Hardware-independent.

Summary and bottom line

The NEW fix is unambiguously beneficial for ReadAccessor<0,1,2> on stencil and coherent point patterns. The gains are 1.5–2.3× and portable across CPU and GPU. ReadAccessor<0> (leaf-only) sees no benefit from the fix and no regression — it is effectively unchanged.

For production use with the NEW accessor, the choice of accessor type matters:

Use case	Recommended accessor
Stencil / convolution / dense sweep	ReadAccessor<0,1,2> (NEW gains 2×)
Sequential / LeafJump access	ReadAccessor<0,1,2> (NEW gains 1.4–2.3×)
NodeJump / random scatter	ReadAccessor<0> (avoids regression, same or faster)
Unknown / mixed	ReadAccessor<0,1,2> — wins on the common workloads

Per-machine takeaways

- LeafJump is the primary CPU win on all machines: 1.65×/1.48× (Desktop), 1.83×/1.40× (Workstation), 1.58×/1.15× (Dell laptop), 1.63× (Razer Laptop) on CPU-1T/GPU for NEW <0,1,2> . This is the pattern that directly exercises the level-1 cache rescue: leaf cache misses but the lower internal node is still cached.
- Stencil GPU results are architecture-dependent. On SM 12.0 (Blackwell): NEW <0,1,2> improves 10 %. On SM 8.9 (Ada): it regresses 12–14 %, confirmed on two independent SM 8.9 machines (Workstation RTX 6000 Ada and Dell laptop RTX 5000 Ada — 0.88× on both). On SM 7.5 (Turing / Razer Laptop): regresses 20 % (0.80×) — the oldest architecture absorbs the extra checks least efficiently. ReadAccessor<0> is neutral in all cases.
- GPU NodeJump wins on Ada and Blackwell; regresses on Turing: 1.32× (Desktop), 1.48× (Workstation), 1.36× (Dell laptop) vs. 0.76× (Razer Laptop). With one representative voxel per actual lower-internal node the level-2 cache hit path pays off on newer GPUs; SM 7.5 Turing cannot absorb the overhead.
- NodeJump CPU-1T regresses for <0,1,2> on the Desktop (0.74×): level-1 always misses with the real topology pattern; the extra cache checks add latency with no benefit. The GPU still wins (1.32×) because the level-2 hit at the upper-node boundary outweighs check overhead.
- Random always regresses for <0,1,2> (0.84–0.94× CPU; 0.79–0.89× GPU across all platforms): all cache levels miss, extra checks are pure overhead. <0> is neutral (0.99–1.02×).

- **ReadAccessor<0>** is unaffected by OLD → NEW on CPU on all platforms and patterns. Its code path (leaf check → root) has no level-1/2 checks to change. On SM 7.5 GPU, <0> also shows no meaningful change (0.81–1.02×).
- **CPU multi-thread scaling:** On the 32-thread 9950X (Desktop), LeafJump achieves ~15× (MT vs 1T), Random ~16×, Sequential ~6×. On the 64-thread Threadripper PRO (Workstation), LeafJump achieves 13–16×, Random 29–31×, Sequential 7×. On the 12-thread Razer Laptop, LeafJump achieves ~7× and the 6-core topology lets level-1/2 cache hits contribute more per thread.

Building & running

```
# Configure with the benchmark (and CUDA, for the GPU variant) enabled
cmake -S . -B build -DNANOVDB_BUILD_BENCHMARK=ON -DNANOVDB_USE_CUDA=ON

# Build all four executables
cmake --build build --target \
    bench_accessor_old bench_accessor_new \
    bench_accessor_cuda_old bench_accessor_cuda_new -j

# Run – each binary prints results for both ReadAccessor<0,1,2> and ReadAccessor<0>
cd build/nanovdb/nanovdb/benchmark
./bench_accessor_old      # CPU, OLD accessor, both <0,1,2> and <0>
./bench_accessor_new      # CPU, NEW accessor, both <0,1,2> and <0>
./bench_accessor_cuda_old # GPU, OLD accessor, both <0,1,2> and <0>
./bench_accessor_cuda_new # GPU, NEW accessor, both <0,1,2> and <0>
```

The OLD vs NEW behaviour is selected at compile time per target (-DNANOVDB_USE_OLD_ACCESSOR vs -DNANOVDB_NO_OLD_ACCESSOR). Both accessor types (ReadAccessor<0,1,2> and ReadAccessor<0>) are run within each binary.

Files

File	Purpose
BenchPatterns.h	Shared access-pattern generation (identical coords for CPU & GPU)
BenchAccessor.cc	CPU benchmark (single- and multi-threaded)
BenchAccessorCuda.cu	GPU (CUDA) benchmark
CMakeLists.txt	Builds the OLD/NEW × CPU/GPU targets